Time Series and Sequence Learning

Modeling time series with stochastic processes

Fredrik Lindsten, Linköping University

A graphical representation

Classical regression applied to time series data.

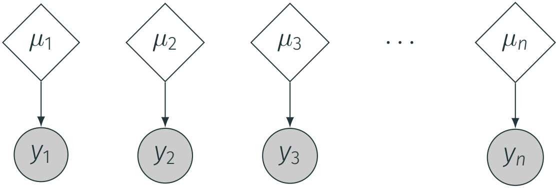

$$\mu_t = \theta_0 + \theta_1 u_t, \qquad u_t = 1993\tfrac{4}{365},\ 1993\tfrac{14}{365},\ \ldots,\ 2020\tfrac{20}{365}$$

Stochastic process

A fundamental approach to time series analysis is to model the data as a stochastic process,

$$\{y_t : t = 1, 2, \ldots\}$$

Stochastic process

A fundamental approach to time series analysis is to model the data as a stochastic process,

$$\{y_t : t = 1, 2, \ldots\}$$

Probabilistic graphical model:

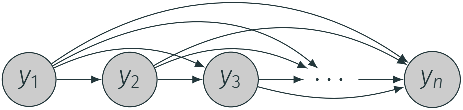

A model for time series data

How can we model the stochastic process?

A complete probabilistic description is given by the *joint probability density function*

$$p(y_{1:n}) = p(y_1, y_2, \ldots, y_n).$$

A model for time series data

How can we model the stochastic process?

A complete probabilistic description is given by the *joint probability density function*

$$p(y_{1:n}) = p(y_1, y_2, \dots, y_n).$$

Cumbersome to work with directly:

1. Dimension grows with time horizon n
2. High-dimensional for large n
3. Can not be used to *forecast* values y_t for $t > n$

A model for time series data

$$p(y_{1:n}) = p(\underline{y_n \mid y_{1:n-1}})\, p(y_{1:n-1})$$

$$= p(y_n \mid y_{1:n-1})\, p(y_{n-1} \mid y_{1:n-2})\, p(y_{1:n-2})$$

$$= p(y_n \mid y_{1:n-1})\, p(y_{n-1} \mid y_{1:n-2}) \cdots p(y_2 \mid y_1)\, \underline{p(y_1)}$$

A model for time series data

Factorize joint pdf as

$$p(y_{1:n}) = \prod_{t=1}^{n} p(y_t \mid y_{1:t-1})$$

It is enough to model the one-step predictive pdf $p(y_t \mid y_{1:t-1})$.

[For notational brevity we define $p(y_1 \mid y_{1:0}) := p(y_1)$]

First-order auto-regressive model

A model for $p(y_t \mid y_{1:t-1})$ tells us how the current value y_t depends on the past values $y_{1:t-1}$.

First-order auto-regressive model

A model for $p(y_t \mid y_{1:t-1})$ tells us how the current value y_t depends on the past values $y_{1:t-1}$.

ex) If we believe that y_t is similar to y_{t-1}, one possibility is to assume

$$y_t = a y_{t-1} + \varepsilon_t, \qquad \varepsilon_t \overset{\text{iid}}{\sim} \mathcal{N}(0, \sigma_\varepsilon^2),$$

for some parameter a.

First-order auto-regressive model

A model for $p(y_t \mid y_{1:t-1})$ tells us how the current value y_t depends on the past values $y_{1:t-1}$.

ex) If we believe that y_t is similar to y_{t-1}, one possibility is to assume

$$y_t = a y_{t-1} + \varepsilon_t, \qquad \varepsilon_t \overset{\text{iid}}{\sim} \mathcal{N}(0, \sigma_\varepsilon^2),$$

for some parameter a.

This is referred to as a (first-order) **auto-regressive (AR) model**.

First-order auto-regressive model

A model for $p(y_t \,|\, y_{1:t-1})$ tells us how the **current value** y_t depends on **the past values** $y_{1:t-1}$.

ex) If we believe that y_t is **similar to** y_{t-1}, one possibility is to assume

$$y_t = a y_{t-1} + \varepsilon_t, \qquad \varepsilon_t \overset{\text{iid}}{\sim} \mathcal{N}(0, \sigma_\varepsilon^2),$$

for some parameter a.

This is referred to as a (first-order) **auto-regressive (AR) model**.

For a first-order AR model we have,

$$p(y_t \,|\, y_{1:t-1}) = p(y_t \,|\, y_{t-1}) = \mathcal{N}\left(y_t \,|\, a y_{t-1}, \sigma_\varepsilon^2\right).$$

First-order auto-regressive model

The dependency structure of the predictive pdf,

$$p(y_t \mid y_{1:t-1}) = p(y_t \mid y_{t-1}) = \mathcal{N}\left(y_t \mid a y_{t-1}, \sigma_\varepsilon^2\right).$$

can be illustrated graphically as

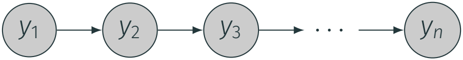

Simulation of AR(1)

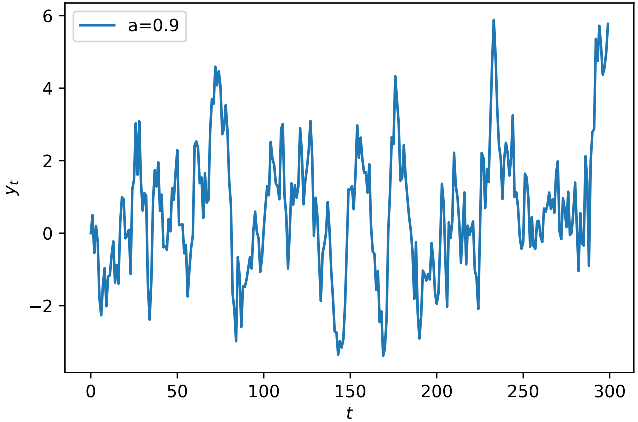

Simulation of AR(1)

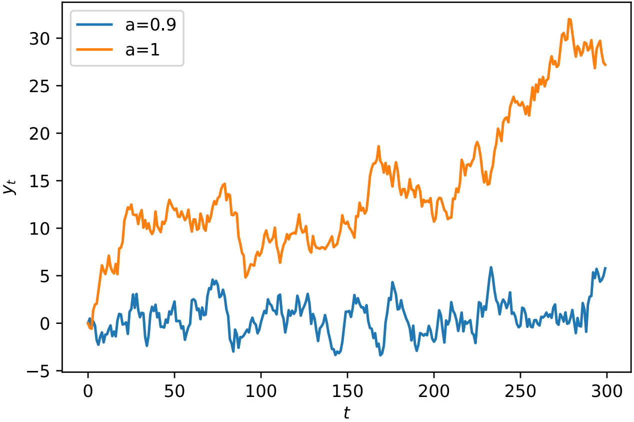

Simulation of AR(1)